

Evaluating and optimizing SQL query with Postgres

Explain Analyze

An HTTP request journey



Get the query and Read the Analyze

What you need:

- A way to get the raw SQL query
- A way to execute the analyze in production and retrieve the result

Getting the raw sql

Django => .query

Sql Alchemy => .str()

SqlModel =>

```
print(query.compile(compile_kwargs={'literal_bindings': True}))
```

- Query log
- Import pdb;
pdb.set_trace()

Live example

An HTTP request journey




Table stats

```
afpy_db=# explain analyze select * from activities;
                                         QUERY PLAN
-----
Seq Scan on activities  (cost=0.00..114787.33 rows=6261033 width=56) (actual time=0.247..163.826 rows=6261121.00 loops=1)
  Buffers: shared hit=282 read=51895
Planning Time: 0.051 ms
Execution Time: 289.072 ms
(4 rows)
```

```
afpy_db=# select relpages, reltuples from pg_class where relname='activities';
   relpages |   reltuples
-----+-----
      52177 | 6.261033e+06
(1 row)
```

```
afpy_db=# select 52177 * 1.0 + 6261033 * 0.01;
           ?column?
-----
          114787.33
(1 row)
```

pg_stats

Ways to optimize for a Query

- 1) Optimize the raw query
- 2) Add an index or remove one
- 3) Upgrade your database spec: `work_mem` and `shared_buffers`

Live example

From sequential scan to index only scan

Conclusion How to Automate the Process?

- HypoPG
- PG analyze

Questions